

Brandon Amos

✉ bda@fb.com • [bamos.github.io](https://github.com/bamos) • [in bdamos](https://www.linkedin.com/in/bdamos) • [brandondamos](https://twitter.com/brandondamos)
🔗 [bamos](https://github.com/bamos) • Last updated on May 27, 2022

Current Position

Research Scientist, *Meta AI (FAIR)*, New York City 2019 – Present

Education

Ph.D. in Computer Science, *Carnegie Mellon University* (0.00/0.00) 2014 – 2019

Thesis: *Differentiable Optimization-Based Modeling for Machine Learning*

Advisor: [J. Zico Kolter](#)

B.S. in Computer Science, *Virginia Tech* (3.99/4.00) 2011 – 2014

Previous Positions

Research Assistant, *Carnegie Mellon University* (with [J. Zico Kolter](#) on ML and optimization) 2016 – 2019

Research Intern, *Intel Labs*, Santa Clara (with [Vladlen Koltun](#) on computer vision) 2018

Research Intern, *Google DeepMind*, London (with [Nando de Freitas](#) and [Misha Denil](#) on RL) 2017

Research Assistant, *Carnegie Mellon University* (with [Mahadev Satyanarayanan](#) on mobile systems) 2014 – 2016

Research Intern, *Adobe Research*, San Jose (with [David Tompkins](#) on distributed systems) 2014

Research Assistant, *Virginia Tech* (with [Layne Watson](#) and [David Easterling](#) on optimization) 2013 – 2014

Research Assistant, *Virginia Tech* (with [Jules White](#) and [Hamilton Turner](#) on mobile systems) 2012 – 2014

Research Assistant, *Virginia Tech* (with [Binoy Ravindran](#) and [Alastair Murray](#) on compilers) 2012 – 2014

Software Intern, *Snowplow*, Remote (Scala development) 2013 – 2014

Software Intern, *Qualcomm*, San Diego (Python and C++ development) 2013

Software Intern, *Phoenix Integration*, Virginia (C++, C#, and Java development) 2012

Network Administrator Intern, *Sunapsys*, Virginia 2011

Honors & Awards

ICLR Outstanding Reviewer 2019

NSF Graduate Research Fellowship 2016 – 2019

Nine undergraduate scholarships 2011 – 2014

Roanoke County Public Schools Engineering, Salem–Roanoke County Chamber of Commerce, Papa John's, Scottish Rite of Freemasonry, VT Intelligence Community Center for Academic Excellence, VT Pamplin Leader, VT Benjamin F. Bock, VT Gay B. Shober, VT I. Luck Gravett

Publications [\[Google Scholar\]](#)

Representative publications that I am a primary author on are **highlighted**.

2022.....

1. *Tutorial on amortized optimization for learning to optimize over continuous domains* [\[code\]](#)

Brandon Amos

arXiv 2022

2. *Cross-Domain Imitation Learning via Optimal Transport* [\[code\]](#)
[Arnaud Fickinger](#), [Samuel Cohen](#), [Stuart Russell](#), and **Brandon Amos**
ICLR 2022

3. *Semi-Discrete Normalizing Flows through Differentiable Tessellation*
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
arXiv 2022

2021

4. *On the model-based stochastic value gradient for continuous reinforcement learning* [code] [slides]
Brandon Amos, Samuel Stanton, Denis Yarats, and Andrew Gordon Wilson
L4DC 2021 (Oral)
5. *Riemannian Convex Potential Maps* [code] [slides]
Samuel Cohen*, **Brandon Amos***, and Yaron Lipman
ICML 2021
6. *CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints* [code]
Anselm Paulus, Michal Rolínek, Vít Musil, **Brandon Amos**, and Georg Martius
ICML 2021
7. *Scalable Online Planning via Reinforcement Learning Fine-Tuning*
Arnaud Fickinger, Hengyuan Hu, **Brandon Amos**, Stuart Russell, and Noam Brown
NeurIPS 2021
8. *Aligning Time Series on Incomparable Spaces* [code] [slides]
Samuel Cohen, Giulia Luise, Alexander Terenin, **Brandon Amos**, and Marc Peter Deisenroth
AISTATS 2021
9. *Learning Neural Event Functions for Ordinary Differential Equations* [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021
10. *Neural Spatio-Temporal Point Processes* [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021
11. *Improving Sample Efficiency in Model-Free Reinforcement Learning from Images* [code]
Denis Yarats, Amy Zhang, Ilya Kostrikov, **Brandon Amos**, Joelle Pineau, and Rob Fergus
AAAI 2021
12. *Neural Fixed-Point Acceleration for Convex Optimization* [code]
Shobha Venkataraman* and **Brandon Amos***
ICML AutoML Workshop 2021
13. *Sliced Multi-Marginal Optimal Transport*
Samuel Cohen, Alexander Terenin, Yannik Pitcan, **Brandon Amos**, Marc Peter Deisenroth, and K S Sesh Kumar
NeurIPS OTML Workshop 2021
14. *Input Convex Gradient Networks*
Jack Richter-Powell, Jonathan Lorraine, and **Brandon Amos**
NeurIPS OTML Workshop 2021
15. *Imitation Learning from Pixel Observations for Continuous Control*
Samuel Cohen, **Brandon Amos**, Marc Peter Deisenroth, Mikael Henaff, Eugene Vinitsky, and Denis Yarats
NeurIPS DeepRL Workshop 2021

16. *MBRL-Lib: A Modular Library for Model-based Reinforcement Learning* [code]
Luis Pineda, **Brandon Amos**, Amy Zhang, Nathan Lambert, and Roberto Calandra
arXiv 2021

2020

17. *The Differentiable Cross-Entropy Method* [code] [slides]
Brandon Amos and Denis Yarats
ICML 2020
18. *Objective Mismatch in Model-based Reinforcement Learning*
Nathan Lambert, **Brandon Amos**, Omry Yadan, and Roberto Calandra
L4DC 2020
19. *QNSTOP: Quasi-Newton Algorithm for Stochastic Optimization* [code]
Brandon Amos, David Easterling, Layne T. Watson, William Thacker, Brent Castle, and Michael Trosset
ACM TOMS 2020
20. *Neural Potts Model*
Tom Sercu, Robert Verkuil, Joshua Meier, **Brandon Amos**, Zeming Lin, Caroline Chen, Jason Liu, Yann LeCun, and Alexander Rives
MLCB 2020
21. *Deep Riemannian Manifold Learning*
Aaron Lou, Maximilian Nickel, and **Brandon Amos**
NeurIPS Geo4dl Workshop 2020

2019

22. *Differentiable Optimization-Based Modeling for Machine Learning* [code]
Brandon Amos
Ph.D. Thesis 2019
23. *Differentiable Convex Optimization Layers* [code]
Akshay Agrawal*, **Brandon Amos***, Shane Barratt*, Stephen Boyd*, Steven Diamond*, and J. Zico Kolter*
NeurIPS 2019
24. *The Limited Multi-Label Projection Layer* [code]
Brandon Amos, Vladlen Koltun, and J. Zico Kolter
arXiv 2019
25. *Generalized Inner Loop Meta-Learning* [code]
Edward Grefenstette, **Brandon Amos**, Denis Yarats, Phu Mon Htut, Artem Molchanov, Franziska Meier, Douwe Kiela, Kyunghyun Cho, and Soumith Chintala
arXiv 2019

2018

26. *Learning Awareness Models*
Brandon Amos, Laurent Dinh, Serkan Cabi, Thomas Rothörl, Sergio Gómez Colmenarejo, Alistair Muldal, Tom Erez, Yuval Tassa, Nando de Freitas, and Misha Denil
ICLR 2018

27. *Differentiable MPC for End-to-end Planning and Control* [code]
Brandon Amos, Ivan Dario Jimenez Rodriguez, Jacob Sacks, Byron Boots, and J. Zico Kolter
NeurIPS 2018
28. *Depth-Limited Solving for Imperfect-Information Games*
Noam Brown, Tuomas Sandholm, and **Brandon Amos**
NeurIPS 2018
29. *Enabling Live Video Analytics with a Scalable and Privacy-Aware Framework*
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM TOMM 2018

2017

30. *OptNet: Differentiable Optimization as a Layer in Neural Networks* [code] [slides]
Brandon Amos and J. Zico Kolter
ICML 2017
31. *Input Convex Neural Networks* [code] [slides]
Brandon Amos, Lei Xu, and J. Zico Kolter
ICML 2017
32. *Task-based End-to-end Model Learning* [code]
Priya L. Donti, **Brandon Amos**, and J. Zico Kolter
NeurIPS 2017
33. *Quasi-Newton Stochastic Optimization Algorithm for Parameter Estimation of a Stochastic Model of the Budding Yeast Cell Cycle*
Minghan Chen, **Brandon Amos**, Layne T. Watson, John Tyson, Yang Cao, Cliff Shaffer, Michael Trosset, Cihan Oguz, and Gisella Kakoti
IEEE/ACM TCBB 2017
34. *You can teach elephants to dance: agile VM handoff for edge computing*
Kiryong Ha, Yoshihisa Abe, Thomas Eiszler, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Rohit Upadhyaya, Padmanabhan Pillai, and Mahadev Satyanarayanan
SEC 2017
35. *An Empirical Study of Latency in an Emerging Class of Edge Computing Applications for Wearable Cognitive Assistance*
Zhuo Chen, Wenlu Hu, Junjue Wang, Siyan Zhao, **Brandon Amos**, Guanhong Wu, Kiryong Ha, Khalid Elgazzar, Padmanabhan Pillai, Roberta Klatzky, Daniel Siewiorek, and Mahadev Satyanarayanan
SEC 2017
36. *A Scalable and Privacy-Aware IoT Service for Live Video Analytics* [code]
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM MMSys 2017 (Best Paper)

2016

37. *OpenFace: A general-purpose face recognition library with mobile applications* [code]
Brandon Amos, Bartosz Ludwiczuk, and Mahadev Satyanarayanan
CMU 2016

38. *Collapsed Variational Inference for Sum-Product Networks*
Han Zhao, Tameem Adel, Geoff Gordon, and **Brandon Amos**
ICML 2016
39. *Quantifying the impact of edge computing on mobile applications*
Wenlu Hu, Ying Gao, Kiryong Ha, Junjue Wang, **Brandon Amos**, Zhuo Chen, Padmanabhan Pillai, and Mahadev Satyanarayanan
ACM SIGOPS 2016
40. *Privacy mediators: helping IoT cross the chasm*
Nigel Davies, Nina Taft, Mahadev Satyanarayanan, Sarah Clinch, and **Brandon Amos**
HotMobile 2016

2015 and earlier.....

41. *Edge Analytics in the Internet of Things*
Mahadev Satyanarayanan, Pieter Simoons, Yu Xiao, Padmanabhan Pillai, Zhuo Chen, Kiryong Ha, Wenlu Hu, and **Brandon Amos**
IEEE Pervasive Computing 2015
42. *Bad Parts: Are Our Manufacturing Systems at Risk of Silent Cyberattacks?*
Hamilton Turner, Jules White, Jaime A. Camelio, Christopher Williams, **Brandon Amos**, and Robert Parker
IEEE Security & Privacy 2015
43. *Early Implementation Experience with Wearable Cognitive Assistance Applications*
Zhuo Chen, Lu Jiang, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai, Alex Hauptmann, and Mahadev Satyanarayanan
WearSys 2015
44. *The Case for Offload Shaping*
Wenlu Hu, **Brandon Amos**, Zhuo Chen, Kiryong Ha, Wolfgang Richter, Padmanabhan Pillai, Benjamin Gilbert, Jan Harkes, and Mahadev Satyanarayanan
HotMobile 2015
45. *Are Cloudlets Necessary?*
Ying Gao, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai, and Mahadev Satyanarayanan
CMU 2015
46. *Adaptive VM handoff across cloudlets*
Kiryong Ha, Yoshihisa Abe, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Padmanabhan Pillai, and Mahadev Satyanarayanan
CMU 2015
47. *Global Parameter Estimation for a Eukaryotic Cell Cycle Model in Systems Biology*
Tricity Andrew, **Brandon Amos**, David Easterling, Cihan Oguz, William Baumann, John Tyson, and Layne T. Watson
SummerSim 2014
48. *Applying machine learning classifiers to dynamic Android malware detection at scale* [code]
Brandon Amos, Hamilton Turner, and Jules White
IWCMC 2013

Repositories

facebookresearch/amortized-optimization-tutorial ★118 <i>Tutorial on amortized optimization</i>	2022
facebookresearch/theseus ★252 <i>Differentiable non-linear optimization library</i>	2022
facebookresearch/rcpm ★56 <i>Riemannian Convex Potential Maps</i>	2021
facebookresearch/svg ★39 <i>Model-based stochastic value gradient</i>	2021
facebookresearch/mbrl-lib ★618 <i>Model-based reinforcement learning library</i>	2021
facebookresearch/dcem ★104 <i>The Differentiable Cross-Entropy Method</i>	2020
facebookresearch/higher ★1.4k <i>PyTorch higher-order gradient and optimization library</i>	2019
bamos/thesis ★268 <i>Ph.D. Thesis LaTeX source code</i>	2019
cvxgrp/cvxpylayers ★1.3k <i>Differentiable Convex Optimization Layers</i>	2019
locuslab/mpc.pytorch ★561 <i>Differentiable Model-Predictive Control</i>	2018
locuslab/icnn ★238 <i>Input Convex Neural Networks</i>	2017
locuslab/optnet ★387 <i>OptNet experiments</i>	2017
locuslab/qpth ★526 <i>Differentiable PyTorch QP solver</i>	2017
bamos/densenet.pytorch ★753 <i>PyTorch DenseNet implementation</i>	2017
bamos/block ★267 <i>Intelligent block matrix constructions</i>	2017
bamos/setGPU ★101 <i>Automatically use the least-loaded GPU</i>	2017
bamos/dcgan-completion.tensorflow ★1.3k <i>Image completion with GANs</i>	2016
cmusatyalab/openface ★14.4k <i>Face recognition with deep neural networks</i>	2015
bamos/zsh-history-analysis ★182 <i>Analyze and plot your zsh history</i>	2014
cparse/cparse ★248 <i>C++ expression parser using Dijkstra's shunting-yard algorithm</i>	2013
bamos/cv ★361 <i>Source for this CV: Creates LaTeX/Markdown from YAML/BibTeX</i>	2013
bamos/dotfiles ★240 <i>Linux, mutt, xmonad, vim, emacs, zsh</i>	2012

Invited Talks

<i>End-to-end model learning for control</i> , ICML Workshop on Decision Awareness in RL	2022
<i>Differentiable optimization-based modeling for machine learning</i> , CPAIOR Master Class	2022
<i>Amortized optimization and learning to optimize</i> , ICCOPT	2022
<i>Modeling and learning paradigms for learning to optimize</i> , SIAM MDS Minisymposium	2022
<i>Learning for control with differentiable optimization and ODEs</i> , Columbia University	2021
<i>Differentiable optimization-based modeling for machine learning</i> , IBM Research	2021
<i>Differentiable optimization for control</i> , Max Planck Institute (Tübingen)	2020
<i>Differentiable optimization-based modeling for machine learning</i> , Mila Seminar	2020
<i>Deep Declarative Networks</i> , ECCV Tutorial	2020
<i>On differentiable optimization for control and vision</i> , CVPR Deep Declarative Networks Workshop	2020
<i>Differentiable optimization-based modeling for machine learning</i> , Caltech CS 159 (Guest Lecture)	2020
<i>Unrolled optimization for learning deep energy models</i> , SIAM MDS Minisymposium	2020
<i>Differentiable optimization-based modeling for machine learning</i> , NYU CILVR Seminar	2019
<i>Differentiable optimization-based modeling for machine learning</i> , INFORMS	2019
<i>Differentiable optimization-based modeling for machine learning</i> , Facebook AI Research	2019
<i>Differentiable optimization-based modeling for machine learning</i> , ISMP	2018
<i>Differentiable optimization-based modeling for machine learning</i> , Google Brain	2018
<i>Differentiable optimization-based modeling for machine learning</i> , Bosch Center for AI	2018
<i>Differentiable optimization-based modeling for machine learning</i> , Waymo Research	2018
<i>Differentiable optimization-based modeling for machine learning</i> , Tesla AI	2018
<i>Differentiable optimization-based modeling for machine learning</i> , NVIDIA Robotics	2018
<i>Differentiable optimization-based modeling for machine learning</i> , Salesforce Research	2018
<i>Differentiable optimization-based modeling for machine learning</i> , OpenAI	2018
<i>Differentiable optimization-based modeling for machine learning</i> , NNAISENSE	2018
<i>Differentiable optimization and control</i> , UC Berkeley	2018

Interns and Students

Aaron Lou (visiting FAIR from Cornell and Stanford)	2020 – 2022
Eugene Vinitzky (visiting FAIR from Berkeley)	2021 – 2022
Arnaud Fickinger (visiting FAIR from Berkeley)	2021 – 2022
Samuel Cohen (visiting FAIR from UCL)	2021 – 2022
Ricky Chen (visiting FAIR from Toronto, now: scientist at FAIR)	2020
Paul Liang (visiting FAIR from CMU)	2020
Phillip Wang (at CMU, now: CEO at Gather)	2018

Professional Activities

NeurIPS Learning Meets Combinatorial Optimization Workshop Organizer	2020
CVPR Deep Declarative Workshop Organizer	2020
ECCV Deep Declarative Tutorial Organizer	2020
CMU CSD MS Admissions	2014 – 2015

Reviewing

Neural Information Processing Systems (NeurIPS)
International Conference on Machine Learning (ICML)
International Conference on Learning Representations (ICLR)
IEEE Conference on Computer Vision and Pattern Recognition (CVPR)
IEEE International Conference on Computer Vision (ICCV)
IEEE International Conference on Robotics and Automation (ICRA)
AAAI Conference on Artificial Intelligence (AAAI)
Optimization Letters

Teaching

Graduate AI (CMU 15-780), TA	S2017
Distributed Systems (CMU 15-440/640), TA	S2016
Software Design and Data Structures (VT CS2114), TA	S2013

Skills

Programming	C, C++, Fortran, Haskell, Java, Lua, Make, Mathematica, Python, R, Scala
Frameworks	JAX, NumPy, Pandas, PyTorch, SciPy, TensorFlow, Torch7
Tools	Linux, emacs, vim, evil, org, mu4e, xmonad, git, tmux, zsh