

Brandon Amos

✉ bda@fb.com • [bamos.github.io](https://github.com/bamos) • [in bdamos](https://www.linkedin.com/in/bdamos) • [brandondamos](https://twitter.com/brandondamos)
[bamos](https://github.com/bamos) • Last updated on October 8, 2021

Current Position

Research Scientist, Facebook AI, New York City 2019 – Present

Education

Ph.D. in Computer Science, Carnegie Mellon University (0.00/0.00) 2014 – 2019

Differentiable Optimization-Based Modeling for Machine Learning

Advisors: J. Zico Kolter (2016 – 2019), Mahadev Satyanarayanan (2014 – 2016)

B.S. in Computer Science, Virginia Tech (3.99/4.00) 2011 – 2014

Research Internships

Intel Labs, Santa Clara (Host: [Vladlen Koltun](#)) 2018

Google DeepMind, London (Hosts: [Misha Denil](#) and [Nando de Freitas](#)) 2017

Adobe Research, San Jose (Host: [David Tompkins](#)) 2014

Honors & Awards

NSF Graduate Research Fellowship 2016 – 2019

Nine undergraduate scholarships 2011 – 2014

Roanoke County Public Schools Engineering, Salem–Roanoke County Chamber of Commerce, Papa John's, Scottish Rite of Freemasonry, VT Intelligence Community Center for Academic Excellence, VT Pamplin Leader, VT Benjamin F. Bock, VT Gay B. Shober, VT I. Luck Gravett

Publications

Representative publications that I am a primary author on are **highlighted**.

[[Google Scholar](#)] [[BibTeX](#)]

2021

1. *On the model-based stochastic value gradient for continuous reinforcement learning* [[code](#)] [[slides](#)]
Brandon Amos, Samuel Stanton, Denis Yarats, and Andrew Gordon Wilson
L4DC 2021 (Oral)
2. *Riemannian Convex Potential Maps* [[code](#)] [[slides](#)]
Samuel Cohen*, **Brandon Amos***, and Yaron Lipman
ICML 2021
3. *CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints* [[code](#)]
Anselm Paulus, Michal Rolínek, Vít Musil, **Brandon Amos**, and Georg Martius
ICML 2021
4. *Scalable Online Planning via Reinforcement Learning Fine-Tuning*
Arnaud Fickinger, Hengyuan Hu, **Brandon Amos**, Stuart Russell, and Noam Brown
NeurIPS 2021
5. *Aligning Time Series on Incomparable Spaces* [[code](#)] [[slides](#)]
Samuel Cohen, Giulia Luise, Alexander Terenin, **Brandon Amos**, and Marc Peter Deisenroth
AISTATS 2021

6. [Learning Neural Event Functions for Ordinary Differential Equations](#) [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021
7. [Neural Spatio-Temporal Point Processes](#) [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021
8. [Improving Sample Efficiency in Model-Free Reinforcement Learning from Images](#) [code]
Denis Yarats, Amy Zhang, Ilya Kostrikov, **Brandon Amos**, Joelle Pineau, and Rob Fergus
AAAI 2021
9. [Neural Fixed-Point Acceleration for Convex Optimization](#) [code]
Shobha Venkataraman* and **Brandon Amos***
ICML AutoML 2021
10. [MBRL-Lib: A Modular Library for Model-based Reinforcement Learning](#) [code]
Luis Pineda, **Brandon Amos**, Amy Zhang, Nathan Lambert, and Roberto Calandra
arXiv 2021
11. [Cross-Domain Imitation Learning via Optimal Transport](#)
Arnaud Fickinger, Samuel Cohen, Stuart Russell, and **Brandon Amos**
arXiv 2021

2020

12. [The Differentiable Cross-Entropy Method](#) [code] [slides]
Brandon Amos and Denis Yarats
ICML 2020
13. [Objective Mismatch in Model-based Reinforcement Learning](#)
Nathan Lambert, **Brandon Amos**, Omry Yadan, and Roberto Calandra
L4DC 2020
14. [QNSTOP: Quasi-Newton Algorithm for Stochastic Optimization](#) [code]
Brandon Amos, David Easterling, Layne T. Watson, William Thacker, Brent Castle, and Michael Trosset
ACM TOMS 2020
15. [Neural Potts Model](#)
Tom Sercu, Robert Verkuil, Joshua Meier, **Brandon Amos**, Zeming Lin, Caroline Chen, Jason Liu, Yann LeCun, and Alexander Rives
MLCB 2020
16. [Deep Riemannian Manifold Learning](#)
Aaron Lou, Maximilian Nickel, and **Brandon Amos**
NeurIPS Geo4dl 2020

2019

17. [Differentiable Optimization-Based Modeling for Machine Learning](#) [code]
Brandon Amos
Ph.D. Thesis 2019
18. [Differentiable Convex Optimization Layers](#) [code]
Akshay Agrawal*, **Brandon Amos***, Shane Barratt*, Stephen Boyd*, Steven Diamond*, and J. Zico Kolter*
NeurIPS 2019

19. *The Limited Multi-Label Projection Layer* [code]
Brandon Amos, Vladlen Koltun, and J. Zico Kolter
arXiv 2019
20. *Generalized Inner Loop Meta-Learning* [code]
Edward Grefenstette, **Brandon Amos**, Denis Yarats, Phu Mon Htut, Artem Molchanov, Franziska Meier, Douwe Kiela, Kyunghyun Cho, and Soumith Chintala
arXiv 2019

2018

21. *Learning Awareness Models*
Brandon Amos, Laurent Dinh, Serkan Cabi, Thomas Rothörl, Sergio Gómez Colmenarejo, Alistair Muldal, Tom Erez, Yuval Tassa, Nando de Freitas, and Misha Denil
ICLR 2018
22. *Differentiable MPC for End-to-end Planning and Control* [code]
Brandon Amos, Ivan Dario Jimenez Rodriguez, Jacob Sacks, Byron Boots, and J. Zico Kolter
NeurIPS 2018
23. *Depth-Limited Solving for Imperfect-Information Games*
Noam Brown, Tuomas Sandholm, and **Brandon Amos**
NeurIPS 2018
24. *Enabling Live Video Analytics with a Scalable and Privacy-Aware Framework*
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM TOMM 2018

2017

25. *OptNet: Differentiable Optimization as a Layer in Neural Networks* [code] [slides]
Brandon Amos and J. Zico Kolter
ICML 2017
26. *Input Convex Neural Networks* [code] [slides]
Brandon Amos, Lei Xu, and J. Zico Kolter
ICML 2017
27. *Task-based End-to-end Model Learning* [code]
Priya L. Donti, **Brandon Amos**, and J. Zico Kolter
NeurIPS 2017
28. *Quasi-Newton Stochastic Optimization Algorithm for Parameter Estimation of a Stochastic Model of the Budding Yeast Cell Cycle*
Minghan Chen, **Brandon Amos**, Layne T. Watson, John Tyson, Yang Cao, Cliff Shaffer, Michael Trosset, Cihan Oguz, and Gisella Kakoti
IEEE/ACM TCBB 2017
29. *You can teach elephants to dance: agile VM handoff for edge computing*
Kiryong Ha, Yoshihisa Abe, Thomas Eiszler, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Rohit Upadhyaya, Padmanabhan Pillai, and Mahadev Satyanarayanan
SEC 2017

30. *An Empirical Study of Latency in an Emerging Class of Edge Computing Applications for Wearable Cognitive Assistance*
Zhuo Chen, Wenlu Hu, Junjue Wang, Siyan Zhao, **Brandon Amos**, Guanhong Wu, Kiryong Ha, Khalid Elgazzar, Padmanabhan Pillai, Roberta Klatzky, Daniel Siewiorek, and Mahadev Satyanarayanan
SEC 2017
31. *A Scalable and Privacy-Aware IoT Service for Live Video Analytics* [code]
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM MMSys 2017 (Best Paper)

2016

32. *OpenFace: A general-purpose face recognition library with mobile applications* [code]
Brandon Amos, Bartosz Ludwiczuk, and Mahadev Satyanarayanan
CMU 2016
33. *Collapsed Variational Inference for Sum-Product Networks*
Han Zhao, Tameem Adel, Geoff Gordon, and **Brandon Amos**
ICML 2016
34. *Quantifying the impact of edge computing on mobile applications*
Wenlu Hu, Ying Gao, Kiryong Ha, Junjue Wang, **Brandon Amos**, Zhuo Chen, Padmanabhan Pillai, and Mahadev Satyanarayanan
ACM SIGOPS 2016
35. *Privacy mediators: helping IoT cross the chasm*
Nigel Davies, Nina Taft, Mahadev Satyanarayanan, Sarah Clinch, and **Brandon Amos**
HotMobile 2016

2015 and earlier

36. *Edge Analytics in the Internet of Things*
Mahadev Satyanarayanan, Pieter Simoens, Yu Xiao, Padmanabhan Pillai, Zhuo Chen, Kiryong Ha, Wenlu Hu, and **Brandon Amos**
IEEE Pervasive Computing 2015
37. *Bad Parts: Are Our Manufacturing Systems at Risk of Silent Cyberattacks?*
Hamilton Turner, Jules White, Jaime A. Camelio, Christopher Williams, **Brandon Amos**, and Robert Parker
IEEE Security & Privacy 2015
38. *Early Implementation Experience with Wearable Cognitive Assistance Applications*
Zhuo Chen, Lu Jiang, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai, Alex Hauptmann, and Mahadev Satyanarayanan
WearSys 2015
39. *The Case for Offload Shaping*
Wenlu Hu, **Brandon Amos**, Zhuo Chen, Kiryong Ha, Wolfgang Richter, Padmanabhan Pillai, Benjamin Gilbert, Jan Harkes, and Mahadev Satyanarayanan
HotMobile 2015
40. *Are Cloudlets Necessary?*
Ying Gao, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai, and Mahadev Satyanarayanan
CMU 2015

41. [Adaptive VM handoff across cloudlets](#)
Kiryong Ha, Yoshihisa Abe, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Padmanabhan Pillai, and Mahadev Satyanarayanan
CMU 2015
42. [Global Parameter Estimation for a Eukaryotic Cell Cycle Model in Systems Biology](#)
Tricity Andrew, **Brandon Amos**, David Easterling, Cihan Oguz, William Baumann, John Tyson, and Layne T. Watson
SummerSim 2014
43. [Applying machine learning classifiers to dynamic Android malware detection at scale](#) [code]
Brandon Amos, Hamilton Turner, and Jules White
IWCMC 2013

Repositories

facebookresearch/mbrl-lib ★478 <i>Model-based reinforcement learning library</i>	2021
facebookresearch/dcem ★92 <i>The Differentiable Cross-Entropy Method</i>	2020
facebookresearch/higher ★1.2k <i>PyTorch higher-order gradient and optimization library</i>	2019
bamos/thesis ★257 <i>Ph.D. Thesis LaTeX source code</i>	2019
cvxgrp/cvxpylayers ★1.1k <i>Differentiable Convex Optimization Layers</i>	2019
locuslab/mpc.pytorch ★483 <i>Differentiable Model-Predictive Control</i>	2018
locuslab/icnn ★229 <i>Input Convex Neural Networks</i>	2017
locuslab/optnet ★371 <i>OptNet experiments</i>	2017
locuslab/qpth ★493 <i>Differentiable PyTorch QP solver</i>	2017
bamos/densenet.pytorch ★713 <i>PyTorch DenseNet implementation</i>	2017
bamos/block ★258 <i>Intelligent block matrix constructions</i>	2017
bamos/setGPU ★97 <i>Automatically use the least-loaded GPU</i>	2017
bamos/dcgan-completion.tensorflow ★1.3k <i>Image completion with GANs</i>	2016
cmusatyalab/openface ★14.2k <i>Face recognition with deep neural networks</i>	2015
bamos/zsh-history-analysis ★162 <i>Analyze and plot your zsh history</i>	2014
bamos/cv ★324 <i>Source for this CV: Creates LaTeX/Markdown from YAML/BibTeX</i>	2013
bamos/dotfiles ★229 <i>Linux, mutt, xmonad, vim, emacs, zsh</i>	2012

Invited Talks

Columbia University	2021
IBM Research	2021
Max Planck Institute for Intelligent Systems (Tübingen) Seminar	2020
Montreal Institute for Learning Algorithms Seminar	2020
ECCV Deep Declarative Networks Tutorial	2020
CVPR Deep Declarative Networks Workshop	2020
Caltech CS 159, Guest Lecture	2020
SIAM MDS Minisymposium on Learning Parameterized Energy Minimization Models	2020
New York University CILVR Seminar	2019
INFORMS Session on Prediction and Optimization	2019
Facebook AI Research	2019
ISMP Session on Machine Learning and Optimization	2018
Google Brain	2018
Bosch Center for AI	2018
Waymo Research	2018
Tesla AI	2018
NVIDIA Robotics	2018
Salesforce Research	2018
OpenAI	2018

NNAISENSE	2018
UC Berkeley	2018

Interns and Students

Samuel Cohen (visiting FAIR from UCL)	2021
Eugene Vinitzky (visiting FAIR from Berkeley)	2021
Arnaud Fickinger (visiting FAIR from Berkeley)	2021
Aaron Lou (visiting FAIR from Cornell, now: Ph.D. student at Stanford)	2020
Ricky Chen (visiting FAIR from Toronto, now: scientist at FAIR)	2020
Paul Liang (visiting FAIR from CMU)	2020
Phillip Wang (at CMU, now: CEO at Gather)	2018

Professional Activities

Reviewing: AAAI, ICML, NeurIPS, ICLR*, ICCV, CVPR, ICRA	*Outstanding reviewer
NeurIPS Learning Meets Combinatorial Optimization Workshop Organizer	2020
CVPR Deep Declarative Workshop Organizer	2020
ECCV Deep Declarative Tutorial Organizer	2020
CMU CSD MS Admissions	2014 – 2015

Teaching

Graduate AI (CMU 15-780), TA	S2017
Distributed Systems (CMU 15-440/640), TA	S2016
Software Design and Data Structures (VT CS2114), TA	S2013

Skills

Programming	C, C++, Fortran, Haskell, Java, Lua, Make, Mathematica, Python, R, Scala
Frameworks	JAX, NumPy, Pandas, PyTorch, SciPy, TensorFlow, Torch7
Tools	Linux, emacs, vim, evil, org, mu4e, xmonad, git, tmux, zsh