Brandon Amos

✉ bda@meta.com ● 🌐 bamos.github.io ● in bdamos ● 🐦 brandondamos
○ bamos ● Last updated on June 13, 2022

Current Position

Research Scientist, *Meta AI (FAIR)*, New York City   2019 – Present

Education

Ph.D. in Computer Science, *Carnegie Mellon University* (0.00/0.00)   2014 – 2019
Thesis: *Differentiable Optimization-Based Modeling for Machine Learning*
Advisor: J. Zico Kolter

B.S. in Computer Science, *Virginia Tech* (3.99/4.00)   2011 – 2014

Previous Positions

Research Assistant, *Carnegie Mellon University* (with J. Zico Kolter on ML and optimization)   2016 – 2019
Research Intern, *Intel Labs*, Santa Clara (with Vladlen Koltun on computer vision)   2018
Research Intern, *Google DeepMind*, London (with Nando de Freitas and Misha Denil on RL)   2017
Research Assistant, *Carnegie Mellon University* (with Mahadev Satyanarayanan on mobile systems)   2014 – 2016
Research Intern, *Adobe Research*, San Jose (with David Tompkins on distributed systems)   2014
Research Assistant, *Virginia Tech* (with Layne Watson and David Easterling on optimization)   2013 – 2014
Research Assistant, *Virginia Tech* (with Jules White and Hamilton Turner on mobile systems)   2012 – 2014
Research Assistant, *Virginia Tech* (with Binoy Ravindran and Alastair Murray on compilers)   2012 – 2014
Software Intern, *Snowplow* (Scala development)   2013 – 2014
Software Intern, *Qualcomm*, San Diego (Python and C++ development)   2013
Software Intern, *Phoenix Integration*, Virginia (C++, C#, and Java development)   2012
Network Administrator Intern, *Sunapsys*, Virginia   2011

Honors & Awards

ICLR Outstanding Reviewer   2019
NSF Graduate Research Fellowship   2016 – 2019
Nine undergraduate scholarships   2011 – 2014
Roanoke County Public Schools Engineering, Salem–Roanoke County Chamber of Commerce, Papa John's, Scottish Rite of Freemasonry, VT Intelligence Community Center for Academic Excellence, VT Pamplin Leader, VT Benjamin F. Bock, VT Gay B. Shober, VT I. Luck Gravett

Publications [Google Scholar; 4841+ citations, h-index: 29+]

Representative publications that I am a primary author on are highlighted.

2022

1. *Tutorial on amortized optimization for learning to optimize over continuous domains* [code]
   Brandon Amos
   arXiv 2022

2. *Cross-Domain Imitation Learning via Optimal Transport* [code]
   Arnaud Fickinger, Samuel Cohen, Stuart Russell, and **Brandon Amos**
   ICLR 2022

3. *Semi-Discrete Normalizing Flows through Differentiable Tessellation*
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
arXiv 2022

4. *Meta Optimal Transport* [code]
Brandon Amos, Samuel Cohen, Giulia Luise, and Ievgen Redko
arXiv 2022

2021

5. *On the model-based stochastic value gradient for continuous reinforcement learning* [code] [slides]
Brandon Amos, Samuel Stanton, Denis Yarats, and Andrew Gordon Wilson
L4DC 2021 (Oral)

6. *Riemannian Convex Potential Maps* [code] [slides]
Samuel Cohen*, **Brandon Amos***, and Yaron Lipman
ICML 2021

7. *CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints* [code]
Anselm Paulus, Michal Rolínek, Vít Musil, **Brandon Amos**, and Georg Martius
ICML 2021

8. *Scalable Online Planning via Reinforcement Learning Fine-Tuning*
Arnaud Fickinger, Hengyuan Hu, **Brandon Amos**, Stuart Russell, and Noam Brown
NeurIPS 2021

9. *Aligning Time Series on Incomparable Spaces* [code] [slides]
Samuel Cohen, Giulia Luise, Alexander Terenin, **Brandon Amos**, and Marc Peter Deisenroth
AISTATS 2021

10. *Learning Neural Event Functions for Ordinary Differential Equations* [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021

11. *Neural Spatio-Temporal Point Processes* [code]
Ricky T. Q. Chen, **Brandon Amos**, and Maximilian Nickel
ICLR 2021

12. *Improving Sample Efficiency in Model-Free Reinforcement Learning from Images* [code]
Denis Yarats, Amy Zhang, Ilya Kostrikov, **Brandon Amos**, Joelle Pineau, and Rob Fergus
AAAI 2021

13. *Neural Fixed-Point Acceleration for Convex Optimization* [code]
Shobha Venkataraman* and **Brandon Amos***
ICML AutoML Workshop 2021

14. *Sliced Multi-Marginal Optimal Transport*
Samuel Cohen, Alexander Terenin, Yannik Pitcan, **Brandon Amos**, Marc Peter Deisenroth, and
K S Sesh Kumar
NeurIPS OTML Workshop 2021

15. *Input Convex Gradient Networks*
Jack Richter-Powell, Jonathan Lorraine, and **Brandon Amos**
NeurIPS OTML Workshop 2021

16. *Imitation Learning from Pixel Observations for Continuous Control*
Samuel Cohen, **Brandon Amos**, Marc Peter Deisenroth, Mikael Henaff, Eugene Vinitsky, and
Denis Yarats
NeurIPS DeepRL Workshop 2021

17. *MBRL-Lib: A Modular Library for Model-based Reinforcement Learning* [code]
Luis Pineda, **Brandon Amos**, Amy Zhang, Nathan Lambert, and Roberto Calandra
arXiv 2021

2020

18. *The Differentiable Cross-Entropy Method* [code] [slides]
Brandon Amos and Denis Yarats
ICML 2020

19. *Objective Mismatch in Model-based Reinforcement Learning*
Nathan Lambert, **Brandon Amos**, Omry Yadan, and Roberto Calandra
L4DC 2020

20. *QNSTOP: Quasi-Newton Algorithm for Stochastic Optimization* [code]
Brandon Amos, David Easterling, Layne T. Watson, William Thacker, Brent Castle, and
Michael Trosset
ACM TOMS 2020

21. *Neural Potts Model*
Tom Sercu, Robert Verkuil, Joshua Meier, **Brandon Amos**, Zeming Lin, Caroline Chen, Jason Liu,
Yann LeCun, and Alexander Rives
MLCB 2020

22. *Deep Riemannian Manifold Learning*
Aaron Lou, Maximilian Nickel, and **Brandon Amos**
NeurIPS Geo4dl Workshop 2020

2019

23. *Differentiable Optimization-Based Modeling for Machine Learning* [code]
Brandon Amos
Ph.D. Thesis 2019

24. *Differentiable Convex Optimization Layers* [code]
Akshay Agrawal*, **Brandon Amos***, Shane Barratt*, Stephen Boyd*, Steven Diamond*, and
J. Zico Kolter*
NeurIPS 2019

25. *The Limited Multi-Label Projection Layer* [code]
Brandon Amos, Vladlen Koltun, and J. Zico Kolter
arXiv 2019

26. *Generalized Inner Loop Meta-Learning* [code]
Edward Grefenstette, **Brandon Amos**, Denis Yarats, Phu Mon Htut, Artem Molchanov,
Franziska Meier, Douwe Kiela, Kyunghyun Cho, and Soumith Chintala
arXiv 2019

2018

27. *Learning Awareness Models*
Brandon Amos, Laurent Dinh, Serkan Cabi, Thomas Rothörl, Sergio Gómez Colmenarejo, Alistair Muldal, Tom Erez, Yuval Tassa, Nando de Freitas, and Misha Denil
ICLR 2018

28. *Differentiable MPC for End-to-end Planning and Control* [code]
Brandon Amos, Ivan Dario Jimenez Rodriguez, Jacob Sacks, Byron Boots, and J. Zico Kolter
NeurIPS 2018

29. *Depth-Limited Solving for Imperfect-Information Games*
Noam Brown, Tuomas Sandholm, and **Brandon Amos**
NeurIPS 2018

30. *Enabling Live Video Analytics with a Scalable and Privacy-Aware Framework*
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM TOMM 2018

2017

31. *OptNet: Differentiable Optimization as a Layer in Neural Networks* [code] [slides]
Brandon Amos and J. Zico Kolter
ICML 2017

32. *Input Convex Neural Networks* [code] [slides]
Brandon Amos, Lei Xu, and J. Zico Kolter
ICML 2017

33. *Task-based End-to-end Model Learning* [code]
Priya L. Donti, **Brandon Amos**, and J. Zico Kolter
NeurIPS 2017

34. *Quasi-Newton Stochastic Optimization Algorithm for Parameter Estimation of a Stochastic Model of the Budding Yeast Cell Cycle*
Minghan Chen, **Brandon Amos**, Layne T. Watson, John Tyson, Yang Cao, Cliff Shaffer, Michael Trosset, Cihan Oguz, and Gisella Kakoti
IEEE/ACM TCBB 2017

35. *You can teach elephants to dance: agile VM handoff for edge computing*
Kiryong Ha, Yoshihisa Abe, Thomas Eiszler, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Rohit Upadhyaya, Padmanabhan Pillai, and Mahadev Satyanarayanan
SEC 2017

36. *An Empirical Study of Latency in an Emerging Class of Edge Computing Applications for Wearable Cognitive Assistance*
Zhuo Chen, Wenlu Hu, Junjue Wang, Siyan Zhao, **Brandon Amos**, Guanhang Wu, Kiryong Ha, Khalid Elgazzar, Padmanabhan Pillai, Roberta Klatzky, Daniel Siewiorek, and Mahadev Satyanarayanan
SEC 2017

37. *A Scalable and Privacy-Aware IoT Service for Live Video Analytics* [code]
Junjue Wang, **Brandon Amos**, Anupam Das, Padmanabhan Pillai, Norman Sadeh, and Mahadev Satyanarayanan
ACM MMSys 2017 (Best Paper)

2016

38. *OpenFace: A general-purpose face recognition library with mobile applications* [code]
Brandon Amos, Bartosz Ludwiczuk, and Mahadev Satyanarayanan
CMU 2016

39. *Collapsed Variational Inference for Sum-Product Networks*
Han Zhao, Tameem Adel, Geoff Gordon, and **Brandon Amos**
ICML 2016

40. *Quantifying the impact of edge computing on mobile applications*
Wenlu Hu, Ying Gao, Kiryong Ha, Junjue Wang, **Brandon Amos**, Zhuo Chen, Padmanabhan Pillai,
and Mahadev Satyanarayanan
ACM SIGOPS 2016

41. *Privacy mediators: helping IoT cross the chasm*
Nigel Davies, Nina Taft, Mahadev Satyanarayanan, Sarah Clinch, and **Brandon Amos**
HotMobile 2016

2015 and earlier

42. *Edge Analytics in the Internet of Things*
Mahadev Satyanarayanan, Pieter Simoens, Yu Xiao, Padmanabhan Pillai, Zhuo Chen, Kiryong Ha,
Wenlu Hu, and **Brandon Amos**
IEEE Pervasive Computing 2015

43. *Bad Parts: Are Our Manufacturing Systems at Risk of Silent Cyberattacks?*
Hamilton Turner, Jules White, Jaime A. Camelio, Christopher Williams, **Brandon Amos**, and
Robert Parker
IEEE Security & Privacy 2015

44. *Early Implementation Experience with Wearable Cognitive Assistance Applications*
Zhuo Chen, Lu Jiang, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai,
Alex Hauptmann, and Mahadev Satyanarayanan
WearSys 2015

45. *The Case for Offload Shaping*
Wenlu Hu, **Brandon Amos**, Zhuo Chen, Kiryong Ha, Wolfgang Richter, Padmanabhan Pillai,
Benjamin Gilbert, Jan Harkes, and Mahadev Satyanarayanan
HotMobile 2015

46. *Are Cloudlets Necessary?*
Ying Gao, Wenlu Hu, Kiryong Ha, **Brandon Amos**, Padmanabhan Pillai, and
Mahadev Satyanarayanan
CMU 2015

47. *Adaptive VM handoff across cloudlets*
Kiryong Ha, Yoshihisa Abe, Zhuo Chen, Wenlu Hu, **Brandon Amos**, Padmanabhan Pillai, and
Mahadev Satyanarayanan
CMU 2015

48. *Global Parameter Estimation for a Eukaryotic Cell Cycle Model in Systems Biology*
Tricity Andrew, **Brandon Amos**, David Easterling, Cihan Oguz, William Baumann, John Tyson, and
Layne T. Watson
SummerSim 2014

49. *Applying machine learning classifiers to dynamic Android malware detection at scale* [code]
    Brandon Amos, Hamilton Turner, and Jules White
    IWCMC 2013

Open Source Repositories

1. facebookresearch/amortized-optimization-tutorial | ★118 | *Tutorial on amortized optimization* 2022
2. facebookresearch/theseus | ★252 | *Differentiable non-linear optimization library* 2022
3. facebookresearch/meta-ot | ★4 | *Meta Optimal Transport* 2022
4. facebookresearch/rcpm | ★56 | *Riemannian Convex Potential Maps* 2021
5. facebookresearch/svg | ★39 | *Model-based stochastic value gradient* 2021
6. facebookresearch/mbrl-lib | ★618 | *Model-based reinforcement learning library* 2021
7. facebookresearch/dcem | ★104 | *The Differentiable Cross-Entropy Method* 2020
8. facebookresearch/higher | ★1.4k | *PyTorch higher-order gradient and optimization library* 2019
9. bamos/thesis | ★268 | *Ph.D. Thesis LaTeX source code* 2019
10. cvxgrp/cvxpylayers | ★1.3k | *Differentiable Convex Optimization Layers* 2019
11. locuslab/mpc.pytorch | ★561 | *Differentiable Model-Predictive Control* 2018
12. locuslab/icnn | ★238 | *Input Convex Neural Networks* 2017
13. locuslab/optnet | ★387 | *OptNet experiments* 2017
14. locuslab/qpth | ★526 | *Differentiable PyTorch QP solver* 2017
15. bamos/densenet.pytorch | ★753 | *PyTorch DenseNet implementation* 2017
16. bamos/block | ★267 | *Intelligent block matrix constructions* 2017
17. bamos/setGPU | ★101 | *Automatically use the least-loaded GPU* 2017
18. bamos/dcgan-completion.tensorflow | ★1.3k | *Image completion with GANs* 2016
19. cmusatyalab/openface | ★14.4k | *Face recognition with deep neural networks* 2015
20. vtopt/qnstop | ★10 | *Fortran package for Quasi-newton stochastic optimization* 2014
21. bamos/snowglobe | ★27 | *Haskell-driven, self-hosted web analytics with minimal configuration* 2014
22. bamos/zsh-history-analysis | ★182 | *Analyze and plot your zsh history* 2014
23. bamos/beamer-snippets | ★106 | *Beamer and TikZ snippets* 2014
24. bamos/latex-templates | ★355 | *LaTeX templates* 2013
25. cparse/cparse | ★248 | *C++ expression parser using Dijkstra's shunting-yard algorithm* 2013
26. bamos/cv | ★361 | *Source for this CV: Creates LaTeX/Markdown from YAML/BibTeX* 2013
27. bamos/python-scripts | ★196 | *Short and fun Python scripts* 2013
28. bamos/reading-list | ★185 | *YAML reading list and notes system* 2013
29. bamos/dotfiles | ★240 | ♥ *Linux, xmonad, emacs, vim, zsh, tmux* 2012

Invited Talks

1. *End-to-end model learning for control*, ICML Workshop on Decision Awareness in RL 2022
2. *Differentiable optimization-based modeling for machine learning*, CPAIOR Master Class 2022
3. *Amortized optimization and learning to optimize*, ICCOPT 2022
4. *Modeling and learning paradigms for learning to optimize*, SIAM MDS Minisymposium 2022
5. *Learning for control with differentiable optimization and ODEs*, Columbia University 2021
6. *Differentiable optimization-based modeling for machine learning*, IBM Research 2021
7. *Differentiable optimization for control*, Max Planck Institute (Tübingen) 2020
8. *Differentiable optimization-based modeling for machine learning*, Mila Seminar 2020
9. *Deep Declarative Networks*, ECCV Tutorial 2020
10. *On differentiable optimization for control and vision*, CVPR Deep Declarative Networks Workshop 2020
11. *Differentiable optimization-based modeling for machine learning*, Caltech CS 159 (Guest Lecture) 2020
12. *Unrolled optimization for learning deep energy models*, SIAM MDS Minisymposium 2020
13. *Differentiable optimization-based modeling for machine learning*, NYU CILVR Seminar 2019

14. *Differentiable optimization-based modeling for machine learning*, INFORMS	2019
15. *Differentiable optimization-based modeling for machine learning*, Facebook AI Research	2019
16. *Differentiable optimization-based modeling for machine learning*, ISMP	2018
17. *Differentiable optimization-based modeling for machine learning*, Google Brain	2018
18. *Differentiable optimization-based modeling for machine learning*, Bosch Center for AI	2018
19. *Differentiable optimization-based modeling for machine learning*, Waymo Research	2018
20. *Differentiable optimization-based modeling for machine learning*, Tesla AI	2018
21. *Differentiable optimization-based modeling for machine learning*, NVIDIA Robotics	2018
22. *Differentiable optimization-based modeling for machine learning*, Salesforce Research	2018
23. *Differentiable optimization-based modeling for machine learning*, OpenAI	2018
24. *Differentiable optimization-based modeling for machine learning*, NNAISENSE	2018
25. *Differentiable optimization and control*, UC Berkeley	2018

Interns and Students

Aaron Lou (visiting FAIR from Cornell and Stanford)	2020 – 2022
Eugene Vinitsky (visiting FAIR from Berkeley)	2021 – 2022
Arnaud Fickinger (visiting FAIR from Berkeley)	2021 – 2022
Samuel Cohen (visiting FAIR from UCL)	2021 – 2022
Ricky Chen (visiting FAIR from Toronto, now: scientist at FAIR)	2020
Paul Liang (visiting FAIR from CMU)	2020
Phillip Wang (at CMU, now: CEO at Gather)	2018

Professional Activities

NeurIPS Learning Meets Combinatorial Optimization Workshop Organizer	2020
CVPR Deep Declarative Networks Workshop Organizer	2020
ECCV Deep Declarative Networks Tutorial Organizer	2020
CMU CSD MS Admissions	2014 – 2015

Reviewing

Neural Information Processing Systems (NeurIPS)
International Conference on Machine Learning (ICML)
International Conference on Learning Representations (ICLR)
IEEE Conference on Computer Vision and Pattern Recognition (CVPR)
IEEE International Conference on Computer Vision (ICCV)
IEEE International Conference on Robotics and Automation (ICRA)
AAAI Conference on Artificial Intelligence
Optimization Letters

Teaching

Graduate AI (CMU 15-780), TA	S2017
Distributed Systems (CMU 15-440/640), TA	S2016
Software Design and Data Structures (VT CS2114), TA	S2013

Skills

Programming	C, C++, Fortran, Haskell, Java, Lua, Make, Mathematica, Python, R, Scala
Frameworks	JAX, NumPy, Pandas, PyTorch, SciPy, TensorFlow, Torch7
Toolbox	Linux, emacs, vim, evil, org, mu4e, xmonad, git, tmux, zsh